

Daily Ticket Spike Detection

Date Initiated: April 2025

Prepared For: Driver Excellence and whomsoever it may concern

Objective

This system identifies **unusual surges (spikes)** in complaint tickets raised by drivers, broken down by **issue category** and **driver zone**, and generates **daily alerts** for proactive intervention.

The aim is to distinguish **normal volume fluctuations** from **true anomalies** based on historical trends, ensuring timely alerts that prompt resolution and root-cause analysis.

Process Flow

1. Data Extraction

- Filters for selecting tickets:
 - Only those tickets that were raised by drivers were considered
 - Last 6 months of data (since October 2024)
 - Only active issue categories and issues were considered
 - Driver's zone was taken for counting tickets
-

2. Aggregation

- Tickets computed are calculated zone-wise and category-wise.
 - Simultaneously, the top-contributing **issue within each category** is used based on the count of tickets and not max change in count of tickets.
-

3. Trend Modeling via Rolling Averages

For each Zone–Category pair:

- Historical daily ticket volumes are analyzed using **rolling windows approach**:
 - **7-day**, **30-day**, and **90-day**
- Metrics computed:
 - Average
 - Standard Deviation
 - Minimum and Maximum

Why Rolling Windows?

They capture short-term (weekly) and long-term (monthly/quarterly) behavior patterns, providing a dynamic, adaptive benchmark for what's "normal."

4. Anomaly Detection Logic

- **Safe Threshold:** Highest of three (**average + standard deviation**) bands
 - Rolling 7 days (average + std deviation)
 - Rolling 30 days (average + std deviation)
 - Rolling 90 days (average + std deviation)
- A **spike** is flagged if:
 - Ticket count exceeds the safe threshold
 - Count based threshold further reduce the number of alerts as:
 - Increase from safe zone should be at least of ≥ 3 tickets count
 - Ticket count of category should be ≥ 5 tickets to be considered

5. Severity Classification

Spikes are categorized as **High**, **Medium**, or **Low** based on:

- **Volume of activity** (High vs Low category historically)
- **Magnitude of increase** (absolute & % increase)

Volume Level	High Severity	Medium Severity	Low Severity
Low Volume (avg < 25)	≥ 20 tickets or $\geq 200\%$	≥ 10 tickets or $\geq 100\%$	<10 tickets or <100%
High Volume (avg ≥ 25)	≥ 50 tickets or $\geq 70\%$	≥ 30 tickets or $\geq 40\%$	<30 tickets or <40%

This adaptive thresholding ensures fair detection across both high-frequency and niche categories.

6. Reporting & Alerts

Google Sheets (Audit Trail)

- **raw:** Replaced daily
- **dailyTicketLogs:** Appended for traceability and checking past data

Slack Alerts

- Posted only for **High** and **Medium** spikes
- Sent to: **#central-alerts** (configurable)
- Message format:
 - Zone, Category
 - Ticket count vs Safe Threshold
 - % increase
 - Top Issue
 - Whether it's the **highest in 7/30/90 days**

What Makes This System Robust

Feature	Benefit
Adaptive windows (7/30/90)	Captures both short-term volatility and seasonal trends
Zone-level granularity	Enables targeted ops intervention
Top issue attribution	Speeds up RCA process
Severity tiers	Filters noise, ensures only action-worthy alerts
Multi-channel outputs	Google sheet + Slack alerts for real-time visibility

Edge Handling & Safeguards




- **Missing trend data:** Handled via historical backfill logic
 - **False positives in low volume areas:** Controlled using thresholds on minimum counts and differences
 - **Slack alert errors:** Fallback logic to avoid complete script failure
 - **Safe zone definition:** Max of 3 bands ensures capturing of seasonal or program based increase in tickets efficiently to decrease chances of false alerts.
-

Slack message meaning

The Slack message includes key visual and textual elements that help teams quickly understand the severity and context of a ticket spike. Below are explanations of specific message components:





1. Peak Note (**peak_note**)

This line adds contextual information about how exceptional the current day's spike is compared to recent history:

-  **Highest in L90!**
→ The current ticket count is the highest in the **last 90 days** — a strong indicator of a rare or serious spike.
-  **Monthly high!**
→ The ticket count is the highest in the **last 30 days**, signaling a recent uptick.
-  **Weekly high!**
→ The count is the highest in the **past week**, flagging short-term urgency.

2. Severity Icons (**severity_emoji**)

Each ticket spike is assigned a severity label and an accompanying **emoji** to make the alert **visually scannable**:

-  – **High Severity**
Indicates a major spike requiring immediate investigation.
-  – **Medium Severity**
Represents a moderate anomaly that should be reviewed soon.
-  – **Low Severity**
A small but valid deviation, flagged for awareness.
-  – **Unknown**
Used when severity couldn't be calculated due to missing or invalid data.

3. Ticket Delta Breakdown ((**+X**, **+Y%**))

This part of the message quantifies the spike using two intuitive metrics:

- **+X** → **Absolute increase** in number of tickets compared to historical average
(e.g., +14 means 14 more tickets than usual)
- **+Y%** → **Percentage increase** relative to average
(e.g., +120% indicates the volume more than doubled)

Both metrics help balance perspective: large percentages in low-volume categories are not overhyped, and small percentages in high-volume zones are not ignored.